

Лабораторная 7

(методические указания)

Задание.

Реализуйте транспонирование матрицы размерностью $N \times K$ без использования разделяемой памяти, с разделяемой памятью без разрешения конфликта банков и с разрешением конфликта банков. Сравните время выполнения соответствующих ядер на GPU. Для всех трёх случаев определите эффективность использования разделяемой памяти с помощью метрик `nvprof` или `psi`.

Цель: приобретение навыков использования разделяемой памяти.

I. Подготовить ядра для тестирования

```
__global__ void glnit(float* a){
    int k=threadIdx.x+blockIdx.x*blockDim.x;
    int n=threadIdx.y+blockIdx.y*blockDim.y;
    int K=blockDim.x*gridDim.x;

    a[k+n*K]=(float)(k+n*K);
}

__global__ void gTranspose1(float* a, float* b){
    int k=threadIdx.x+blockIdx.x*blockDim.x;
    int n=threadIdx.y+blockIdx.y*blockDim.y;
    int K=blockDim.x*gridDim.x;
    int N=blockDim.y*gridDim.y;

    b[k+n*K]=a[n+k*N];
}
```

```
__global__ void gTranspose2(float* a, float* b){  
    int k=threadIdx.x+blockIdx.x*blockDim.x;  
    int n=threadIdx.y+blockIdx.y*blockDim.y;  
    int K=blockDim.x*gridDim.x;  
    int N=blockDim.y*gridDim.y;  
  
    b[n+k*N]=a[k+n*K];  
}
```

```
#define SH_DIM 32
__global__ void gTransposeSM(float* a, float* b){
    __shared__ float cache[SH_DIM][SH_DIM];
    int k=threadIdx.x+blockIdx.x*blockDim.x;
    int n=threadIdx.y+blockIdx.y*blockDim.y;
    int N=blockDim.x*gridDim.x;

    cache[threadIdx.y][threadIdx.x]=a[k+n*N];
    __syncthreads();

    k=threadIdx.x+blockIdx.y*blockDim.x;
    n=threadIdx.y+blockIdx.x*blockDim.y;
    b[k+n*N]=cache[threadIdx.x][threadIdx.y];
}
```

```
__global__ void gTransposeSM_WC(float* a, float* b){  
    __shared__ float cache[SH_DIM][SH_DIM+1];  
    int k=threadIdx.x+blockIdx.x*blockDim.x;  
    int n=threadIdx.y+blockIdx.y*blockDim.y;  
    int N=blockDim.x*gridDim.x;  
  
    cache[threadIdx.y][threadIdx.x]=a[k+n*N];  
    __syncthreads();  
  
    k=threadIdx.x+blockIdx.y*blockDim.x;  
    n=threadIdx.y+blockIdx.x*blockDim.y;  
    b[k+n*N]=cache[threadIdx.x][threadIdx.y];  
}
```

II. Написать драйвер для тестирования

```
#include <stdio.h>
#define CUDA_CHECK_RETURN(value) {\
    cudaError_t  _m_cudaStat = value;\
    if (_m_cudaStat != cudaSuccess) {\
        fprintf(stderr, "Error %s at line %d in file %s\n",\
            cudaGetErrorString(_m_cudaStat), __LINE__, __FILE__);\
        exit(1);\
    }\
}
```

```
int main(int argc, char* argv){
    if(argc<3){
        fprintf(stderr, "USAGE: tr_mat-25 <dimension of matrix> <dimension of
                                                                    threads>\n");
        return -1;
    }
    int N=atoi(argv[1]);
    int dim_of_threads=atoi(argv[2]);
    if(N%dim_of_threads){
        fprintf(stderr, "change dimensions\n");
        return -1;
    }
    int dim_of_blocks=N/dim_of_threads;
    const int max_size=1<<8;
    if(dim_of_blocks>max_size){
        fprintf(stderr, "too many blocks\n");
        return -1;
    }
}
```

```
float *a, *b;
```

```
cudaMalloc((void**)&a, N*N*sizeof(float));
```

```
cudaMalloc((void**)&b, N*N*sizeof(float));
```

```
glnit<<<dim3(dim_of_blocks, dim_of_blocks),  
         dim3(dim_of_threads, dim_of_threads)>>>(a);
```

```
cudaDeviceSynchronize();
```

```
CUDA_CHECK_RETURN(cudaGetLastError());
```

```
cudaMemset(b, 0, N*N*sizeof(float));
```



```
gTranspose1<<<dim3(dim_of_blocks, dim_of_blocks),  
             dim3(dim_of_threads, dim_of_threads)>>>(a,b);  
cudaDeviceSynchronize(); CUDA_CHECK_RETURN(cudaGetLastError());  
gTranspose2<<<dim3(dim_of_blocks, dim_of_blocks),  
             dim3(dim_of_threads, dim_of_threads)>>>(a,b);  
cudaDeviceSynchronize(); CUDA_CHECK_RETURN(cudaGetLastError());  
gTransposeSM<<<dim3(dim_of_blocks, dim_of_blocks),  
              dim3(dim_of_threads, dim_of_threads)>>>(a,b);  
cudaDeviceSynchronize(); CUDA_CHECK_RETURN(cudaGetLastError());  
gTransposeSM_WC<<<dim3(dim_of_blocks, dim_of_blocks),  
               dim3(dim_of_threads, dim_of_threads)>>>(a,b);  
cudaDeviceSynchronize(); CUDA_CHECK_RETURN(cudaGetLastError());  
  
cudaFree(a);  
cudaFree(b);  
}
```

III. Подобрать метрики для анализа производительности

NSIGHT COMPUTE COMMAND LINE INTERFACE

v2023.3.1 | October 2023

User Manual

6.3. Metric Comparison

nvprof Metric	PerfWorks Metric or Formula (\geq SM 7.0)
shared_efficiency	smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct
shared_load_throughput	l1tex__data_pipe_lsu_wavefronts_mem_shared_op_ld.sum.per_second

III. Провести профилирование

```
Lab6> nvprof ./tr_mat25 256 32
```

```
==9499== NVPROF is profiling process 9499, command: ./tr_mat25 256 32
```

```
==9499== Profiling application: ./tr_mat25 256 32
```

```
==9499== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	29.34%	9.9530us	1	9.9530us	9.9530us	9.9530us	
gTranspose2(float*, float*)							
	22.36%	7.5840us	1	7.5840us	7.5840us	7.5840us	
gTransposeSM(float*, float*)							
	21.22%	7.1990us	1	7.1990us	7.1990us	7.1990us	
gTranspose1(float*, float*)							
	11.04%	3.7440us	1	3.7440us	3.7440us	3.7440us	
gInit(float*)							
	9.91%	3.3600us	1	3.3600us	3.3600us	3.3600us	
gTransposeSM_WC(float*, float*)							

```
...> nvprof -m shared_efficiency ./tr_mat25 256 32
```

Invocations	Metric Name		Metric Description
	Min	Max	Avg
Device "GeForce GTX 560 Ti (0)"			
Kernel: gTranspose1(float*, float*)			
1	shared_efficiency		Shared Memory Efficiency
	0.00%	0.00%	0.00%
Kernel: gTranspose2(float*, float*)			
	0.00%	0.00%	0.00%
Kernel: gTransposeSM(float*, float*)			
	6.06%	6.06%	6.06%
Kernel: gTransposeSM_WC(float*, float*)			
	100.00%	100.00%	100.00%

```
/Lab6> ncu --metrics
```

```
smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct ./tr_mat25 256 32
```

```
gTranspose1(float *, float *) (8, 8, 1)x(32, 32, 1), Context 1, Stream 7, Device 0, CC 7.5
```

```
Section: Command line profiler metrics
```

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct	%	0

```
gTranspose2(float *, float *) (8, 8, 1)x(32, 32, 1), Context 1, Stream 7, Device 0, CC 7.5
```

```
Section: Command line profiler metrics
```

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct	%	0

gTransposeSM(float *, float *) (8, 8, 1)x(32, 32, 1), Context 1, Stream 7, Device 0, CC 7.5
Section: Command line profiler metrics

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct	%	6,06

gTransposeSM_WC(float *, float *) (8, 8, 1)x(32, 32, 1), Context 1, Stream 7, Device 0, CC 7.5
Section: Command line profiler metrics

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_wavefront_mem_shared.pct	%	100